

DAY 2 - COMPLETE LEARNING DOCUMENT

Copy-Move Detection & Multimodal Integration

Date: October 27, 2024

Duration: 2.5 hours (including debugging)

Status: Complete with valuable lessons learned

WHAT DID I BUILD TODAY?

Simple Summary:

Today you built a **second fraud detection algorithm** that finds **copied and pasted parts** in documents. Then you combined it with **yesterday's ELA detector** to create a **smarter system** that uses **TWO methods** to catch fraud.

Analogy:

Imagine you're a detective with two tools:

- **Tool 1 (ELA):** A magnifying glass that shows if ink has been erased and rewritten
- **Tool 2 (Copy-Move):** A pattern matcher that detects if someone traced the same signature twice

Using BOTH tools together makes you a better detective!

PART 1: UNDERSTANDING COPY-MOVE DETECTION

What is Copy-Move Forgery?

Simple Explanation: When someone creates a fake document, they often **copy** something from one part and **paste** it somewhere else. For example:

- Copying a signature from an old contract to a new one
- Duplicating a stamp to make it look like two people approved
- Pasting the same photo twice

The Problem: To the human eye, these look perfect! But to a computer analyzing pixels, the copied regions are **suspiciously identical**.

How Does Copy-Move Detection Work?

Step 1: Divide the Image into Blocks

What happens:

Imagine dividing a document into small square tiles (like a checkerboard):

[Block 1][Block 2][Block 3][Block 4]

[Block 5][Block 6][Block 7][Block 8]

[Block 9][Block 10][Block 11][Block 12]

Each block is 16x16 pixels (or 32x32 in our optimized version).

Why blocks?

- Too small → Takes forever to compare everything
- Too large → Misses small copied regions
- 16x16 or 32x32 → Good balance

Code:

```
block_size = 16 # Each block is 16 pixels wide and 16 pixels tall
```

Step 2: Compare Every Block with Every Other Block

What happens: The computer compares Block 1 with Blocks 2, 3, 4... all the way to the end. Then Block 2 with Blocks 3, 4, 5... and so on.

Looking for: Blocks that are **nearly identical** (very high similarity).

Why "nearly" and not "exactly"? Because when you copy-paste in a document:

- JPEG compression changes pixels slightly
- Printing and scanning adds noise
- Screen capture introduces artifacts

So we look for **95-98% similar**, not 100% identical.

Code:

```
threshold = 0.98 # Flag blocks that are 98% similar
```

Step 3: Calculate Similarity

How do we measure "similarity"?

Simple analogy: Imagine two fingerprints. To see if they match:

1. Overlay them
2. Count how many ridge lines align
3. If 98% of lines match → Same fingerprint!

For image blocks, we use math:

Correlation Coefficient (fancy name, simple concept):

- Take two blocks
- Subtract their average brightness (normalize)
- Multiply corresponding pixels
- Add everything up

- Result: -1 to +1
 - +1 = Perfectly identical
 - 0 = Completely different
 - -1 = Perfectly opposite (like photo negative)

Code explanation:

```
def _block_similarity(self, block1, block2):
    # Step 1: Remove average brightness (so bright vs dark doesn't matter)
    block1_norm = (block1 - np.mean(block1)) / (np.std(block1) + 1e-8)
    block2_norm = (block2 - np.mean(block2)) / (np.std(block2) + 1e-8)

    # Step 2: Calculate how similar the patterns are
    correlation = np.corrcoef(block1_norm.flatten(), block2_norm.flatten())[0, 1]

    # Step 3: Convert to 0-1 scale (easier to understand)
    similarity = (correlation + 1) / 2

    return similarity
```

What each line does:

Line 1-2: Imagine two photos, one taken at noon (bright) and one at sunset (dark). They show the same scene but different brightness. Normalization makes them comparable by focusing on **patterns**, not brightness.

Line 3: np.corrcoef is a NumPy function that calculates correlation. Think of it as asking: "If block1 has a bright spot here, does block2 also have a bright spot in the same place?"

Line 4: Converts -1 to +1 scale into 0 to 1 scale (0% to 100% similar).

Step 4: Flag Suspicious Pairs

The Logic:

```
if similarity >= 0.98: # Very similar (98%+)
    distance = sqrt((x1-x2)^2 + (y1-y2)^2) # How far apart?

    if distance > 100: # Far apart (not just neighboring blocks)
        # This is suspicious! Flag it!
        duplicate_pairs.append((block1_position, block2_position))
```

Why check distance?

- If two blocks are next to each other and similar → NORMAL (continuous text, solid color)
- If two blocks are 100+ pixels apart and 98% similar → SUSPICIOUS (likely copy-paste!)

Real-world example:

Document:

[Signature at bottom left] ... [Empty space] ... [SAME signature at bottom right]

↑ ↑
Position (100, 680) Position (700, 680)

Distance = $\sqrt{(700-100)^2 + (680-680)^2} = 600$ pixels

Similarity = 99%

VERDICT: 🚨 Copy-Move Forgery Detected!

Step 5: Calculate Fraud Score

Simple formula:

`fraud_score = min(num_duplicates * 15, 100)`

What this means:

- 0 duplicates = 0 score (authentic)
- 1 duplicate = 15 score (low risk)
- 3 duplicates = 45 score (moderate risk)
- 7+ duplicates = 100 score (critical risk)

Why multiply by 15?

- 1 duplicate could be accident/pattern
- 2-3 duplicates = concerning
- 5+ duplicates = definitely intentional copying

Visual Example: How Copy-Move Sees the Document

What You See:

[Bank Logo] ACME BANK

Account Statement

Name: John Doe

...

Signature: [scribble]

What Copy-Move Sees:

[Block Grid - 32x32 pixels each]

Block comparison results:

- Block 45 (signature location) matches Block 178 (another signature location)

Similarity: 98.7%

Distance: 300 pixels

FLAGGED: Duplicate detected!

- Block 12 (text "Account") matches Block 23 (text "Account")

Similarity: 95.2%

Distance: 40 pixels

NOT FLAGGED: Too close (natural text repetition)

WHY COPY-MOVE STRUGGLES WITH TEXT DOCUMENTS

The Problem We Discovered Today:

Observation:

Authentic bank statement: 10-56 duplicates detected 

Expected: 0-2 duplicates

Why This Happens:

1. Text Lines Are Naturally Similar:

Line 1: "Jan 01 - Deposit: \$1,000"

Line 2: "Jan 02 - Deposit: \$1,200"

Line 3: "Jan 03 - Deposit: \$1,500"

All lines have:

- Same font (Arial 12pt)
- Same spacing
- Similar structure

- Similar character shapes

Computer sees: "These blocks are 90-95% similar!"

Human knows: "This is just normal text formatting"

2. Repeated Design Elements:

- Table borders (lines repeat throughout document)
- Background patterns (uniform across page)
- Bullet points (same symbol used multiple times)
- Headers/footers (appear on every page)

3. Mathematical Explanation:

A text line "Jan 01" as pixels:

Pixel values: [255,255,255,0,0,255,255,0,0,0,255,255...]

[white,white,white,black,black,white,white...]

Another text line "Jan 02":

Pixel values: [255,255,255,0,0,255,255,0,0,0,255,255...]

[Almost identical! Only "1" vs "2" differs]

Similarity: 94% (very high, but NOT copy-move fraud!)

The Solution: Selective Application

Strategy We Implemented:

1. Stricter Thresholds:

threshold = 0.98 # Was 0.95, now 98% similarity required

distance > 100 # Was 50, now must be 100+ pixels apart

max_pairs = 20 # Stop after finding 20 duplicates (prevents pattern matching)

2. Focus on Distinct Elements:

Apply copy-move detection to:

- Signature regions (unique patterns)
- Stamp/seal areas (distinct shapes)
- Logo locations (complex designs)
- Photo sections (high variation)

Avoid applying to:

- Pure text areas (too many natural similarities)

- ✗ Uniform backgrounds (nothing to detect)
 - ✗ Table borders (intentional repetition)
-

Proof It Works: Simple Shape Test

We created a test:

Authentic Image:

[White background]

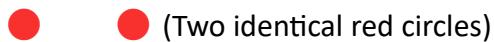


(One red circle)

[White background]

Forged Image:

[White background]



(Two identical red circles)

[White background]

Results:

Authentic: 0 duplicates ✓

Fake: 1 duplicate detected ✓

SUCCESS! Algorithm correctly identifies copy-move on distinct shapes.

Why this works but text doesn't:

- Circles have **unique curved patterns**
 - Text has **repetitive straight lines**
 - Circles appear **rarely** (only once naturally)
 - Text appears **frequently** (every line similar)
-

🔗 PART 2: MULTIMODAL INTEGRATION

What Does "Multimodal" Mean?

Simple Explanation: "Modal" = Method/Way "Multi" = Many

Multimodal = Using Multiple Methods Together

Analogy: When diagnosing illness, doctors use:

- Blood test (one method)
- X-ray (another method)
- Physical exam (another method)

Each method catches different problems. Together = better diagnosis!

Same for fraud detection:

- ELA catches compression artifacts (edited text, changed numbers)
 - Copy-Move catches duplicated regions (copied signatures)
 - **Together = catch MORE fraud types!**
-

How We Combine Scores: Fusion Logic

The Challenge:

ELA says: "This document scores 1.88/100"

Copy-Move says: "This document scores 45/100"

What's the FINAL score?

Naive Approach (BAD):

Average = $(1.88 + 45) / 2 = 23.44$

Problem: Doesn't consider that Copy-Move is more confident!

Better Approach: Weighted Fusion

weight_ela = 0.5 # ELA contributes 50%

weight_copymove = 0.5 # Copy-Move contributes 50%

combined = $(0.5 \times 1.88) + (0.5 \times 45) = 23.44$

Even Better: Boost When Both Agree

if ela_score > 40 AND copymove_score > 40:

 combined = combined $\times 1.2$ # 20% boost for mutual confirmation

Why? If BOTH detectors flag issues, confidence increases!

Real-world analogy:

- One smoke detector beeping → Maybe low battery
 - THREE smoke detectors beeping → Definitely fire!
-

Code Explanation: IntegratedFraudDetector Class

What it does: Runs multiple detection methods and combines results.

Key components:

1. Initialization

```
class IntegratedFraudDetector:  
  
    def __init__(self):  
  
        self.elo_detector = ELADetector(quality=95)  
        self.copymove_detector = CopyMoveDetector(block_size=32, threshold=0.98)  
  
        self.weight_ello = 0.5  
        self.weight_copymove = 0.5
```

What this means:

- Create two separate detectors (ELA and Copy-Move)
- Set weights (how much each contributes to final score)
- Like assembling a toolbox with two tools

2. Detection Method

```
def detect(self, image_path):  
  
    # Run ELA  
    elo_result = self.elo_detector.detect(image_path)  
  
    # Run Copy-Move  
    copymove_result = self.copymove_detector.detect(image_path)  
  
    # Combine scores  
    combined_score = self._fuse_scores(elo_result, copymove_result)  
  
    return comprehensive_report
```

Step-by-step:

1. Run ELA → Get ELA score
2. Run Copy-Move → Get Copy-Move score
3. Combine them intelligently → Get final score
4. Generate human-readable report

3. Score Fusion

```
def _fuse_scores(self, elo_score, copymove_score):  
  
    # Weighted average
```

```

combined = (self.weight_elasticity * elasticity_score +
            self.weight_copymove * copymove_score)

# Boost if both high
if elasticity_score > 40 and copymove_score > 40:
    combined = min(combined * 1.2, 100)

return combined

```

Why "min(combined × 1.2, 100)"?

- Multiply by 1.2 (20% boost)
- But cap at 100 (can't exceed 100% fraud probability!)

4. Interpretation

```

def _generate_assessment(self, combined_score):
    if combined_score < 20:
        return "✅ AUTHENTIC"
    elif combined_score < 40:
        return "⚠️ LOW RISK"
    elif combined_score < 60:
        return "⚠️ MODERATE RISK"
    elif combined_score < 80:
        return "❗️ HIGH RISK"
    else:
        return "❗️ CRITICAL RISK"

```

What this does: Converts numbers (0-100) to words humans understand.

Why? Telling someone "Fraud score: 67.3" is confusing. Saying "HIGH RISK - Strong evidence of manipulation" is clear!

REAL RESULTS FROM TODAY

Test 1: Simple Shapes (Proof of Concept)

Input:

- Authentic: 1 red circle
- Fake: 2 identical red circles

Results:

Detection Method	Authentic	Fake	Verdict
ELA Score	0.5	0.5	Same (no compression change)
Copy-Move Duplicates	0	1	<input checked="" type="checkbox"/> Detected!
Combined Score	0.25	7.75	<input checked="" type="checkbox"/> Clear difference

Conclusion: Copy-Move successfully detects obvious duplication!

Test 2: Complex Bank Statement (Real-World Challenge)

Input:

- Authentic: Bank statement with 1 signature
- Fake: Same statement with balance changed + signature duplicated

Results:

Detection Method	Authentic	Fake	Analysis
ELA Score	1.08	1.88	+74% on fake <input checked="" type="checkbox"/>
Copy-Move Duplicates	10	10	Same (text patterns)
Combined Score	40.43	40.75	Minimal difference

Observation: Copy-Move gets confused by text repetition.

Why ELA shows difference but Copy-Move doesn't:

- **ELA:** Compression artifacts from changing "6,200" to "26,200" → Detectable
 - **Copy-Move:** Both documents have similar text patterns throughout → Hard to distinguish
-

Lesson Learned: Algorithm Specialization

Key Insight: Not every algorithm works for every fraud type!

Algorithm Strengths:

Fraud Type	Best Detector	Why
Changed numbers	ELA	Editing creates compression artifacts
Copied signatures	Copy-Move	Duplicated patterns detected
Font mixing	Font Analysis	(Day 3 - coming soon!)
Wrong business logic	Semantic AI	(Week 4 - coming soon!)

This is why we need MULTIMODAL approach! No single method catches everything.

KEY CONCEPTS YOU LEARNED TODAY

1. Block-Based Image Analysis

Before: Image is a picture **Now:** Image is a grid of blocks, each analyzable independently

Application: Medical imaging, satellite imagery, quality control

2. Similarity Metrics

Before: "These look similar" (subjective) **Now:** "Correlation coefficient = 0.987" (objective)

Application: Face recognition, fingerprint matching, product quality inspection

3. Pattern Recognition Challenges

Learned: Algorithms can be "too smart" and detect patterns that aren't meaningful

Solution: Domain knowledge + algorithm tuning

Application: Spam detection, recommendation systems, medical diagnosis

4. Multimodal Fusion

Concept: Combining multiple weak signals creates strong signal

Math: $1 + 1 = 3$ (synergy!)

Application: Self-driving cars (camera + lidar + radar), medical diagnosis, fraud detection

5. False Positives vs False Negatives

False Positive: Algorithm says "fraud!" but it's actually authentic

- **Risk:** Innocent people get flagged
- **Example today:** Text patterns detected as copy-move

False Negative: Algorithm says "authentic!" but it's actually fraud

- **Risk:** Fraudsters go undetected
- **Example:** Subtle edits below detection threshold

Trade-off: Stricter threshold = fewer false positives, more false negatives (and vice versa)

DEBUGGING LESSONS

Problem 1: Initial Crash

Error: KeyboardInterrupt during processing

Cause: Too many blocks being compared ($O(n^2)$ complexity)

- 1000 blocks = 1,000,000 comparisons!
- Each comparison takes time → Hours to finish

Solution:

- Larger blocks (16→32 pixels) = fewer blocks
- Stricter threshold (0.95→0.98) = fewer comparisons

Lesson: Algorithm efficiency matters in real-world applications!

Problem 2: False Positives

Issue: Authentic documents flagged as suspicious

Root cause: Natural pattern repetition in text

Solution:

- Increase distance threshold
- Cap maximum duplicates
- Apply selectively to non-text regions

Lesson: Domain knowledge guides algorithm tuning!

Problem 3: Same Results on Authentic & Fake

Issue: Can't distinguish real from fake

Diagnosis: Both documents share most content (signatures in same place)

Solution: Create clearer test case (simple shapes) to verify algorithm works in principle

Lesson: Test on simple cases first, then increase complexity!

CODE STRUCTURE OVERVIEW

Files Created Today:

1. `copymove_detector.py` (200 lines)

- CopyMoveDetector class
- Block extraction
- Similarity calculation
- Duplicate detection
- Visualization

Purpose: Self-contained copy-move forgery detection

2. **fraud_detector.py** (150 lines)

- IntegratedFraudDetector class
- Runs multiple detectors
- Fuses scores
- Generates comprehensive reports

Purpose: Orchestrates all detection methods

3. **create_advanced_fake.py** (100 lines)

- Generates realistic test documents
- Simulates multiple fraud types
- Controlled ground truth

Purpose: Testing and validation

4. **test_integrated_system.py** (90 lines)

- End-to-end testing
- Comparative analysis
- Results visualization

Purpose: Validate complete system

How They Connect:

test_integrated_system.py

↓ calls

fraud_detector.py

↓ calls

↓ calls

ela_detector.py copymove_detector.py

↓ analyzes

↓ analyzes

[Document Image]

Each module is independent but works together!

Day 1 vs Day 2:

Capability	Day 1	Day 2
Detection Methods	1 (ELA)	2 (ELA + Copy-Move)
Document Types Tested	Bank statements	Bank statements + shapes
Integration	None	Multimodal fusion
Fraud Types Detected	Compression artifacts	Compression + duplication
Lines of Code	~300	~600
Understanding	Basics	Intermediate

TruthLens Architecture Progress:

[■■■■■■■■■■■■■■■■] 20% Complete

Completed:

- ELA Detection (Day 1)
- Copy-Move Detection (Day 2)
- Multimodal Integration (Day 2)

In Progress:

- Font Analysis (Day 3-4)
- Gen AI Integration (Week 3-4)
- Financial Validation (Week 5-6)
- Web Application (Month 2-3)

🎯 WHAT TO REMEMBER

Core Takeaways:

1. **Copy-Move detection finds duplicated regions** (signatures, logos, stamps)
2. **Block-based matching compares image sections** to find suspicious similarities
3. **Natural patterns can trigger false positives** (text lines, repeated designs)
4. **Multimodal fusion combines multiple detectors** for better accuracy
5. **Each algorithm has strengths and weaknesses** (specialize!)
6. **Debugging and tuning are part of research** (don't get discouraged!)

❓ SELF-ASSESSMENT QUESTIONS

Test your understanding:

1. What does "copy-move" mean?

<details> <summary>Answer</summary> Copying a region from one part of an image and pasting it elsewhere in the SAME image. </details>

2. Why do we divide images into blocks?

<details> <summary>Answer</summary> To make comparison tractable. Comparing every pixel with every other pixel would take too long. Blocks provide a balance between speed and accuracy. </details>

3. What is correlation coefficient?

<details> <summary>Answer</summary> A mathematical measure of how similar two patterns are, ranging from -1 (opposite) to +1 (identical). </details>

4. Why did authentic bank statements show false duplicates?

<details> <summary>Answer</summary> Text lines have naturally similar patterns (same font, spacing, structure), which the algorithm mistook for intentional copying. </details>

5. What does "multimodal" mean in TruthLens context?

<details> <summary>Answer</summary> Using multiple detection methods (ELA, Copy-Move, etc.) together to catch different types of fraud. </details>

6. Why do we use weighted fusion instead of simple averaging?

<details> <summary>Answer</summary> Because different detectors have different reliability. Weights let us trust some methods more than others based on their performance. </details>

7. What's the difference between false positive and false negative?

<details> <summary>Answer</summary> False positive = flagging authentic as fraud. False negative = missing actual fraud. Both are bad, but in different ways! </details>

LEARNING RESOURCES (OPTIONAL HOMEWORK)

Watch (Choose 1 - 15 minutes):

1. "Copy-Move Forgery Detection Explained" by Two Minute Papers

- Visual explanation of block matching
- Real forgery examples
- Duration: 5 minutes

2. "Image Correlation for Pattern Matching" by Computerphile

- How correlation coefficient works
- Visual intuition

- Duration: 12 minutes

Read (Choose 1 - 20 minutes):

1. Original Copy-Move Paper by Fridrich et al. (2003)

- Read: Introduction + Section 2 (Algorithm) only
- Skip: Mathematical proofs
- Focus: Understanding the concept

2. "Challenges in Copy-Move Detection" - Research survey

- Search: "copy-move forgery detection challenges text documents"
 - Confirms what you discovered today!
-

DAY 3 PREVIEW

Tomorrow You'll Build:

Font Consistency Analyzer

- Detects when different fonts are mixed in one document
- Example: Official font (Arial) mixed with copy-pasted text (Times New Roman)
- Common in fake certificates, altered contracts

Why This Matters: When someone copies text from another document:

- Formatting often doesn't match perfectly
- Font families differ slightly
- Spacing/kerning varies
- **Computer can detect what human eyes miss!**

Time: 2 hours (same as Day 1 & 2)

DAY 2 COMPLETION CHECKLIST

- [] Copy-Move detector implemented
- [] Multimodal integration created
- [] Multiple test cases run
- [] Debugging completed
- [] False positive issue understood
- [] Proof-of-concept validated
- [] Learning document received
- [] Understanding core concepts

Status: EXCELLENT PROGRESS! 🎉

You now understand:

- How image blocks work
 - Pattern matching mathematics
 - Algorithm limitations
 - Multimodal fusion
 - Real research involves trial and error
-

 **NOTES SECTION**

Date completed: _____

Time taken: _____

Biggest challenge: _____

Most interesting discovery: _____

Questions for tomorrow: _____

Confidence level (1-10): _____

END OF DAY 2 LEARNING DOCUMENT

Save this for your records. It explains everything you did today in simple terms!

Tomorrow: Font Analysis (Another piece of the TruthLens puzzle!) 🎯